

BUSINESS REQUIREMENTS DOCUMENT

Enterprise Code Knowledge Platform

Business Requirements Document

Field	Detail
Document type	Business Requirements Document (BRD)
Version	1.1
Date	4 September 2026
Audience	Product owners, architects, delivery teams, sponsors and business stakeholders
Purpose	Define what the platform must do, why it matters, and how success will be demonstrated.

Summary: The platform turns a large, complicated software estate into a searchable, evidence-based knowledge base. Engineers, business users and AI assistants can ask how the software works, what a change may affect, where a rule is enforced, and how a business process flows — with answers traceable to the underlying software.

Contents

- 1.**Executive Summary
- 2.**Business Problem
- 3.**Goals and Success Measures
- 4.**Users and Their Needs
- 5.**Scope
- 6.**Platform Capabilities — Simple View
- 7.**Detailed Business Requirements (CAP-1 to CAP-9)
- 8.**Non-Functional Requirements
- 9.**Cross-Cutting Business Rules
- 10.**Data Requirements
- 11.**Delivery Plan
- 12.**Dependencies
- 13.**Key Risks and Controls
- 14.**Acceptance — How We Know It Works
- 15.**Worked Business Scenarios
- 16.**Requirement Summary

1. Executive Summary

Large organisations run software that has been built and changed over many years, by many people, most of whom have left. The knowledge of how that software actually works is scattered across the source code itself, a few long-serving staff, and out-of-date documents. Nobody holds the full picture.

This creates a daily cost. A simple question — *“if we change this rule, what else breaks?”* — takes days to answer, and the answer is often wrong. New joiners take months to become productive. Changes are made cautiously and slowly because the blast radius is unknown.

At the same time, AI assistants have become genuinely useful to engineers, but only when they are given accurate context. Pointed at a large, old, multi-system estate with no context, an assistant guesses. Confident guesses about business logic are worse than no answer at all.

The Enterprise Code Knowledge Platform solves both problems with one asset. It reads the organisation’s own software, and turns it into a structured, queryable, always-refreshable knowledge base that describes the estate in three connected layers:

1. **Structure** — what exists: the components, entry points, data stores, integration points, and how they call each other. Derived automatically from the software itself, so it is never out of date by more than one refresh.
2. **Meaning** — what it is for: the business capabilities, end-to-end processes, business rules, business entities and the people affected by them. Curated by humans, but every statement is anchored to a specific place in the real software with recorded evidence.
3. **Behaviour** — how work actually flows: the sequence of steps in a process, what is validated where, what data gets written, what can fail, and what guards or switches control it.

The platform then exposes this knowledge through a set of question-shaped services that both people and AI assistants can call — over an open, standards-based interface — so an assistant answers from the organisation’s verified knowledge instead of from guesswork.

Safety principle: The platform is read-only and non-invasive. It does not modify the systems it describes.

2. Business Problem

2.1 Current situation

- Different applications are built using different approaches and contain different parts of the business journey.
- A single business process may cross several applications.
- Documentation may describe intended behaviour rather than current behaviour.
- Important knowledge is concentrated in a small number of experienced people.
- AI assistants lack enough structured context to reliably understand a large legacy estate.

2.2 Business impact

ID	Business impact	Priority
P1	Change impact is unknown. Teams may over-test or under-test because they cannot reliably see everything affected.	High
P2	Knowledge is concentrated in individuals, creating bottlenecks and knowledge-loss risk.	High
P3	New engineers spend too long building a mental model of the estate.	High
P4	Production investigation requires manual tracing across applications.	High
P5	Business and engineering use different languages for the same rules and processes.	Medium
P6	AI can produce confident but incorrect answers without verified context.	High
P7	Audit/compliance questions require repeated manual investigation.	Medium

3. Goals and Success Measures

ID	Objective	Type
OBJ-1	Make change impact predictable — reviewed impact list in under 1 hour.	Target
OBJ-2	Reduce dependence on named experts — 80%+ of routine how-does-this-work questions answered without consulting a named expert.	Target
OBJ-3	Shorten onboarding — at least 50% reduction in time for a new engineer to make an accepted change.	Target
OBJ-4	Speed investigation — at least 60% reduction in time from issue report to responsible logic.	Target
OBJ-5	Make AI assistance trustworthy — 90%+ correct and every answer traceable.	Target
OBJ-6	Create one shared business/engineering view — one agreed description per in-scope process.	Target
OBJ-7	Keep knowledge current — refreshable on demand and no older than one agreed refresh cycle.	Target
OBJ-8	Zero operational risk — zero changes required to existing live systems.	Target

4. Users and Their Needs

- **Engineers:** understand code, dependencies, impact and behaviour faster.
- **New joiners:** learn the estate without an expert beside them.
- **Architects / technical leads:** assess change impact and application relationships.
- **Business analysts / product owners:** understand how business rules are enforced.
- **Support / operations:** move from symptoms to responsible logic quickly.
- **QA:** identify what a change touches and plan testing.
- **Compliance / audit / risk:** obtain repeatable, evidenced answers.
- **Subject matter experts:** record knowledge once instead of repeatedly explaining it.

- **Platform owner:** manage scope, refresh, quality and access.

5. Scope

5.1 Included

- Registering enterprise software assets in scope.
- Automatically deriving software structure and relationships.
- Adding human-curated business meaning backed by software evidence.
- Describing end-to-end processes, including cross-application handovers.
- Keeping periodic read-only snapshots of approved configuration/reference values.
- Providing question-focused answer services.
- Providing an open interface for AI assistants and other tools.
- Refreshing knowledge and reporting freshness.
- Access control, audit logging and read-only guarantees.

5.2 Out of scope for this release

- Changing, fixing or refactoring described software.
- Automatically generating or applying code changes.
- Live real-time connections to production systems.
- Customer, personal or transactional business data.
- Replacing source control, issue tracking or documentation tools.
- Runtime monitoring or alerting.
- Automated test execution.
- Describing third-party software the enterprise cannot inspect.


6. Platform Capabilities — Simple View

ID	Capability	Priority
CAP-1	Estate Registration — defines which software is covered.	Must
CAP-2	Structural Knowledge — automatically understands what exists and how it connects.	Must
CAP-3	Business Meaning — add business-language meaning backed by evidence.	Must
CAP-4	Process & Behaviour — explain end-to-end flow, checks, effects and failures.	Must
CAP-5	Meaning-Based Retrieval — find relevant knowledge using ordinary language.	Must
CAP-6	Reference Data — keep approved read-only configuration snapshots.	Must
CAP-7	Answer Services — provide focused question-answering services.	Must
CAP-8	Assistant Interface — let AI assistants discover and use services.	Must

ID	Capability	Priority
CAP-9	Refresh, Trust & Governance — keep knowledge current, traceable and controlled.	Must

Capability Dependencies

Enterprise Code Knowledge Platform · CAP-1 – CAP-9



7. Detailed Business Requirements

The requirement IDs below are retained from the source BRD; the wording is simplified for easier reading.

CAP-1 — Estate Registration

ID	Requirement	Priority
BR-01	Maintain one clear register of all software assets in scope, including stable ID, business name, role, owner and source location.	Must
BR-02	Adding, removing or changing an asset's scope must require only a register change.	Must
BR-03	Support assets built using different technical approaches.	Must

ID	Requirement	Priority
BR-04	Allow exclusions such as generated/third-party content and record the reason.	Should
BR-05	Let users see exactly what is in scope.	Must

CAP-2 — Structural Knowledge

ID	Requirement	Priority
BR-06	Automatically derive each registered asset's structure from the asset itself; no manual entry required.	Must
BR-07	Identify components, operations, entry points, stores and exchange/integration points.	Must
BR-08	Record relationships such as invokes, reads, writes and belongs-to.	Must
BR-09	Point every derived item to its exact source location.	Must
BR-10	Label indirect cross-application connections as inferred and show supporting evidence.	Must
BR-11	Processing an unchanged source twice must produce the same result.	Must
BR-12	Define allowed item/relationship types centrally; unknown types must fail visibly.	Must
BR-13	Report interpretation coverage and what could not be interpreted.	Must

CAP-3 — Business Meaning

ID	Requirement	Priority
BR-14	Store business meaning separately from technical structure, in business language.	Must
BR-15	Support capabilities, processes, rules, entities and affected stakeholders.	Must
BR-16	Every business statement must link to specific software locations.	Must
BR-17	Record evidence and confidence for every anchor.	Must
BR-18	If an anchor no longer resolves, refresh must fail clearly rather than serve stale meaning.	Must
BR-19	Business-literate users must be able to author/review meaning with version history.	Must
BR-20	Clearly distinguish derived facts from curated interpretation in answers.	Must
BR-21	Suggest candidate meaning for human review; never publish unreviewed suggestions as fact.	Should

CAP-4 — Process and Behaviour

ID	Requirement	Priority
BR-22	Describe a business process as ordered stages from entry to completion.	Must

ID	Requirement	Priority
BR-23	Show cross-application handovers within the same end-to-end process.	Must
BR-24	Show all known entry points when there is more than one.	Must
BR-25	List checks/validations in order and where they are enforced.	Must
BR-26	List what information a process records or changes and where.	Must
BR-27	Identify configuration flags/switches that change behaviour and explain them.	Should
BR-28	Identify resource reservation/locking points and explain the business reason.	Should
BR-29	List failure paths and where each failure originates.	Must
BR-30	Describe calculations and their input dependencies.	Should
BR-31	Identify where a new check should be added in the process and software.	Should

CAP-5 — Meaning-Based Retrieval

ID	Requirement	Priority
BR-32	Allow ordinary business-language questions without requiring technical names.	Must
BR-33	Retrieve by meaning, not only exact words.	Must
BR-34	Search business meaning and software content, targeting either or both.	Must
BR-35	Include exact source location with every result.	Must
BR-36	Allow narrowing by asset, process or capability.	Should
BR-37	Refresh incrementally so unchanged material is not reprocessed.	Should
BR-38	Return the full content of a specific source location on request.	Must

CAP-6 — Reference Data

ID	Requirement	Priority
BR-39	Keep an approved copy of configuration/reference values that affect processes.	Must
BR-40	Take periodic snapshots; no live operational connection.	Must
BR-41	Maintain one explicit, reviewable list of copied configuration items.	Must
BR-42	Do not copy customer, personal or transactional data.	Must
BR-43	Make the reference-data copy read-only by design with no write path.	Must
BR-44	Show snapshot age.	Must
BR-45	If queried by users, allow read-only retrieval and validate requests.	Must
BR-46	Answer configuration-driven questions directly.	Must

CAP-7 — Answer Services

ID	Requirement	Priority
BR-47	Provide individually callable services, each answering a clear question category.	Must
BR-48	Cover navigation, impact, flow, checks, effects, explanation, placement, search, detail, configuration and status.	Must
BR-49	Provide composite services for change impact, failure tracing, input acceptance and process description.	Must
BR-50	Composite services return evidence plus guidance for interpreting/presenting it.	Must
BR-51	Define each service in one authoritative place.	Must
BR-52	Each service explains when to use it and how it differs from similar services.	Must
BR-53	All services are read-only.	Must
BR-54	Every answer includes verifiable source references.	Must
BR-55	If the platform does not know, it must say so and explain what is needed; never guess.	Must
BR-56	Services should accept business identifiers, not only technical IDs.	Should
BR-57	Provide a controlled read-only advanced query facility for experts.	Should

CAP-8 — Assistant Interface

ID	Requirement	Priority
BR-58	Expose services through an open, vendor-neutral interface for compliant AI assistants.	Must
BR-59	Support local and shared-network usage.	Must
BR-60	Authenticate shared requests and identify caller/client.	Must
BR-61	Guide assistants to answer only from platform-supplied knowledge.	Must
BR-62	Provide a human-readable guide covering scope, limitations and good questions.	Should
BR-63	Make services discoverable to newly connected assistants.	Must
BR-64	New services automatically become available in every access mode.	Must

CAP-9 — Refresh, Trust & Governance

ID	Requirement	Priority
BR-65	Rebuild complete knowledge from registered assets on demand without manual steps.	Must
BR-66	Make rebuilds safe and repeatable; no corruption or duplication.	Must
BR-67	Run rebuilds in stages and allow restart from a chosen stage.	Should
BR-68	Rebuild one knowledge area without destroying unrelated knowledge.	Must
BR-69	Report freshness and knowledge size/coverage.	Must

ID	Requirement	Priority
BR-70	Fail clearly on inconsistencies instead of silently omitting knowledge.	Must
BR-71	Control access and log caller, time and request.	Must
BR-72	State scope, interpretation coverage and known gaps.	Must
BR-73	Assign a named owner accountable for scope, refresh, quality and access.	Must
BR-74	Define/publish refresh cadence and visibly flag stale knowledge.	Should

8. Non-Functional Requirements

ID	Requirement	Priority
NFR-1	Routine questions should take a few seconds; composite questions roughly 15 seconds or less.	Must
NFR-2	Initial full rebuild should complete in under about 100 minutes.	Must
NFR-3	Support at least several hundred thousand described items and relationships.	Must
NFR-4	Doubling registered assets must not require platform redesign.	Must
NFR-5	Sampled review should show at least 90% factual correctness; every answer verifiable.	Must
NFR-6	Shared instances should be available during business hours; short outages acceptable.	Should
NFR-7	Access authenticated and logged; stored information stays within defined sensitivity boundaries.	Must
NFR-8	New users can ask useful questions in ordinary language without special syntax training.	Must
NFR-9	Unfamiliar teams can add assets/services using documented steps.	Must
NFR-10	Deploy inside the enterprise boundary without sending source outside it.	Must
NFR-11	Ongoing cost should be measurable and justifiable against engineering time saved.	Should

9. Cross-Cutting Business Rules

- **Read-only everywhere:** never write to systems the platform reads from.
- **Everything is traceable:** every fact points to software evidence or a named, dated statement.
- **No silent gaps:** missing knowledge becomes a visible failure or explicit “not known.”
- **Facts vs interpretation:** users can always distinguish them.
- **Non-invasive:** described systems are not modified or instrumented.
- **Configuration-based extension:** adding assets does not require redesign.
- **Single definition:** shared concepts have one authoritative definition.
- **No sensitive data:** customer, personal and transactional data are excluded.

10. Data Requirements

- Generated knowledge is stored separately from source assets.
- Generated output is never written back into source assets.
- Human-curated content is separate, version-controlled and backed by author/date/reason.
- Every knowledge item records provenance, creation time and derived/curated status.
- Credentials and secrets are not stored with knowledge or in source control.
- Complete knowledge is reproducible from registered assets plus curated content.
- Generated knowledge is disposable/rebuildable; curated knowledge is retained and backed up.

11. Delivery Plan

Phase	Scope	Business value
Phase 1	Foundation — CAP-1, CAP-2, navigation/detail. Makes the estate visible and traceable.	Onboarding and orientation
Phase 2	Meaning & Access — CAP-3, CAP-5, CAP-8. Makes the estate understandable and AI-usable.	Shared understanding + AI grounding
Phase 3	Behaviour & Impact — CAP-4 and impact/flow/check/effect/placement services.	Safer changes + faster investigation
Phase 4	Configuration & Composite Answers — CAP-6 and composite services.	End-to-end realistic answers
Phase 5	Trust & Operation — CAP-9.	Freshness, governance and sustainability

12. Dependencies

- Read access to all in-scope source assets.
- Business/engineering SME time for meaning review.
- Data-governance approval for configuration snapshots.
- Named platform owner.
- Enterprise environment for the shared instance.
- Compatible AI assistant access for assistant-grounding objectives.
- Privacy approval for configuration snapshot contents.

13. Key Risks and Controls

ID	Risk	Control
R-1	Wrong answers destroy trust.	Trace answers; say “not known”; never guess.
R-2	Knowledge becomes stale.	Refresh, freshness reporting and broken-anchor validation.
R-3	Business meaning is never curated.	Named owner, SME time and human-reviewed suggestions.
R-4	Scope grows uncontrollably.	Explicit register and phased delivery.
R-5	Engineers perceive surveillance.	Position as an assistive tool; be transparent about logging.
R-6	Some assets cannot be fully interpreted.	Coverage reporting and explicit gaps.
R-7	Configuration snapshot contains sensitive data.	Explicit copy list, exclusions and privacy review.
R-8	Inferred cross-application relationships are wrong.	Label inference and show evidence.
R-9	Platform becomes a single point of failure.	Keep it assistive, not operational.
R-10	Refresh becomes too slow.	Time budget, incremental refresh and restartable stages.

14. Acceptance — How We Know It Works

- **Impact:** a real historical change is analysed; at least 90% of actual impact is found with sources.
- **Onboarding:** a person unfamiliar with the estate explains an end-to-end process using only the platform in under one hour.
- **Investigation:** a real historical issue is traced to responsible logic in under five minutes.
- **AI grounding:** ten realistic questions are answered; at least nine are correct and all are traceable.
- **Currency:** after a source change and refresh, answers reflect it; before refresh, staleness is reported.
- **Safety:** audit confirms no write path to source systems and no sensitive customer/personal/transactional data.

15. Worked Business Scenarios

Scenario 1 — Changing an approval rule

Today: engineers search multiple applications and consult experienced colleagues. With the platform: one question returns where the threshold is used, affected processes, dependent checks, configuration values, stakeholders and evidence.

Scenario 2 — New engineer learning a process

The engineer asks how the process works and receives ordered stages, checks, application handovers, failure paths and source links, without needing an expert's calendar.

Scenario 3 — Investigating a rejected operation

The reported failure is traced to its origin, triggering condition, process stage, relevant configuration and source location.

Scenario 4 — Audit question

The platform returns the business rule, enforcement locations, supporting evidence and confidence.

Scenario 5 — Adding a new application

Register the application and refresh. It becomes describable and its integration points can appear without platform redesign.

16. Requirement Summary

The source BRD contains 74 business requirements across nine capabilities, plus 8 cross-cutting rules, 11 non-functional requirements and 7 data requirements. This refined edition keeps the same requirement IDs and overall intent while simplifying the language and making the business flow easier to understand.

Item	Count
Business requirements (BR-01 to BR-74)	74
Capabilities (CAP-1 to CAP-9)	9
Cross-cutting business rules	8
Non-functional requirements	11
Data requirements	7
Business impact statements (P1 to P7)	7
Objectives (OBJ-1 to OBJ-8)	8
Key risks (R-1 to R-10)	10